

Memory Allocator Discussion

5/9/22: The Curtain Behind malloc()

A Brief Motivation

There are 3 “locations” data for a program can reside in: .data segment, stack, and heap.

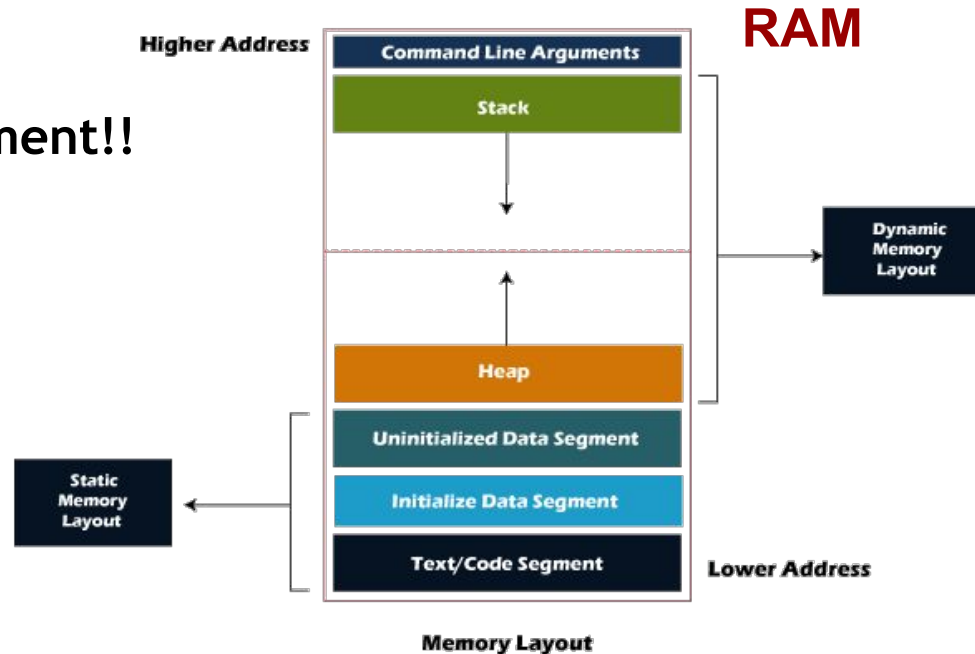
Example Code:

```
unsigned int fifo = 0;    ← .data segment!!
```

```
void loop() {  
    char input[128];      ← stack!!
```

```
    ...
```

```
}
```



Spaces are allocated on the heap with an explicit “malloc” call.

Example Code:

```
unsigned int fifo = 0;    ← .data segment!!
```

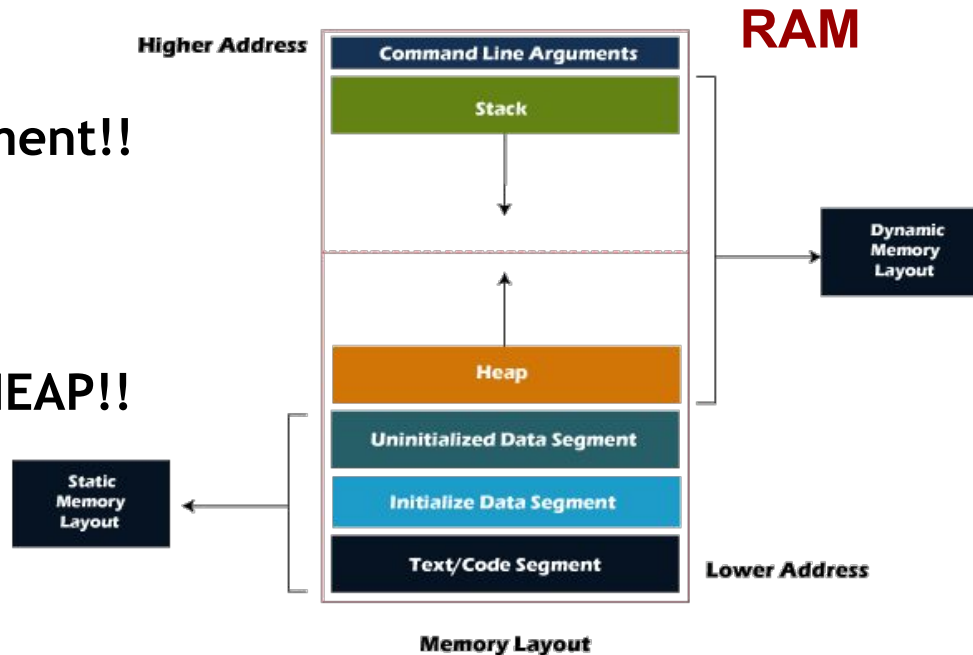
```
void loop() {
```

```
    char input[128];      ← stack!!
```

```
    char *ex = (char*) malloc(4); ← HEAP!!
```

```
    ...
```

```
}
```



Spaces allocated in .data segment and stack must be determinable at compile-time.

Example Code:

`unsigned int fifo = 0;` ← **4 bytes allocated!**

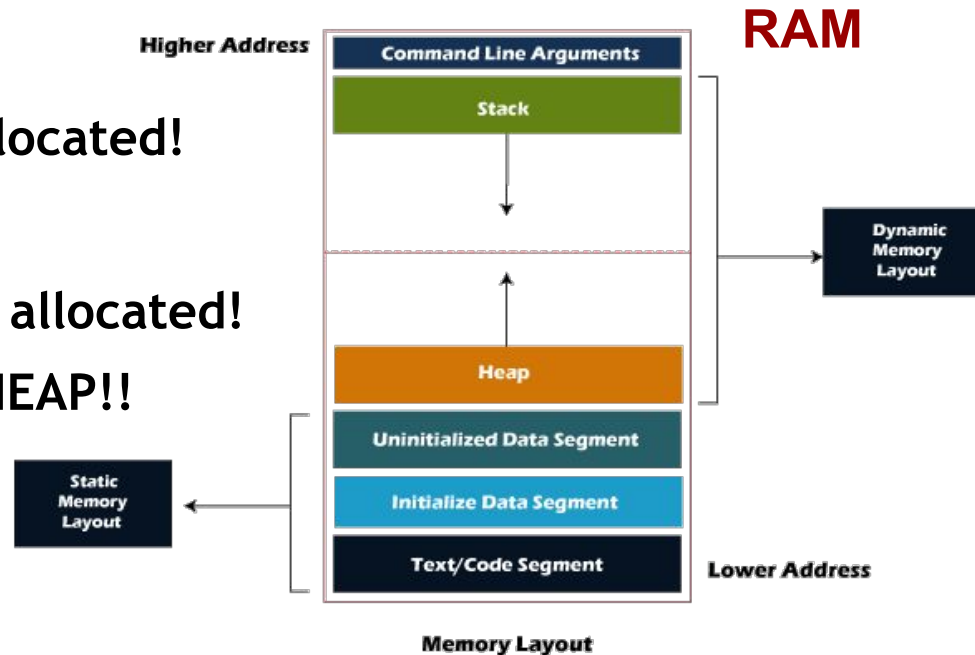
`void loop() {`

`char input[128];` ← **128 bytes allocated!**

`char *ex = (char*) malloc(4);` ← **HEAP!!**

`...`

`}`



Spaces allocated in .data segment and stack must be determinable at compile-time.

Example Code:

`unsigned int fifo = 0;` ← **4 bytes allocated!**

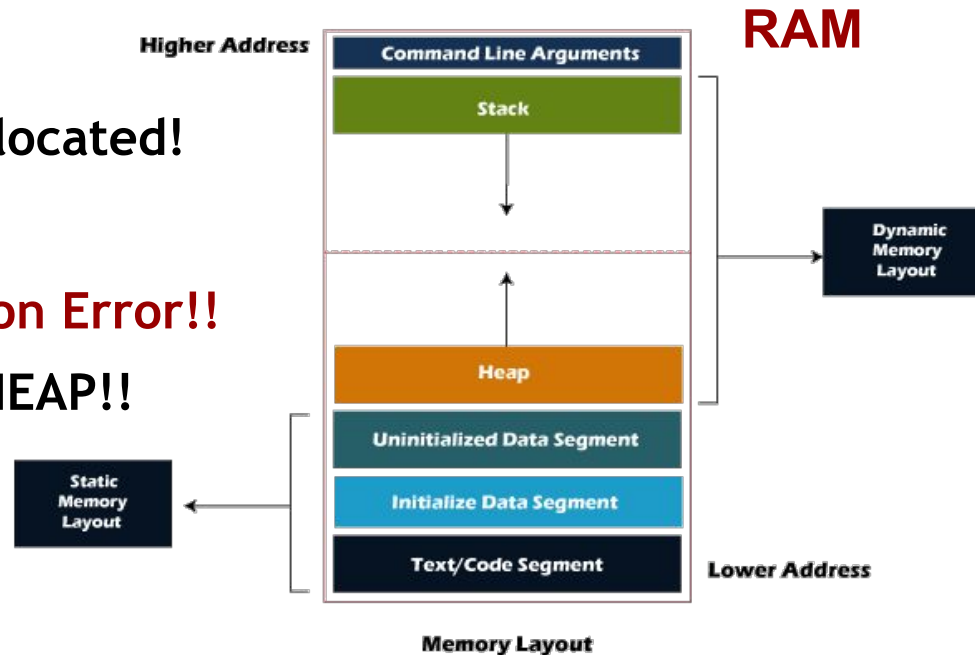
`void loop(int size) {`

`char input[size];` ← **Compilation Error!!**

`char *ex = (char*) malloc(4);` ← **HEAP!!**

`...`

`}`



Spaces allocated on the heap can be determined at run-time instead.

Example Code:

```
unsigned int fifo = 0;    ← 4 bytes allocated!
```

```
void loop(int size) {
```

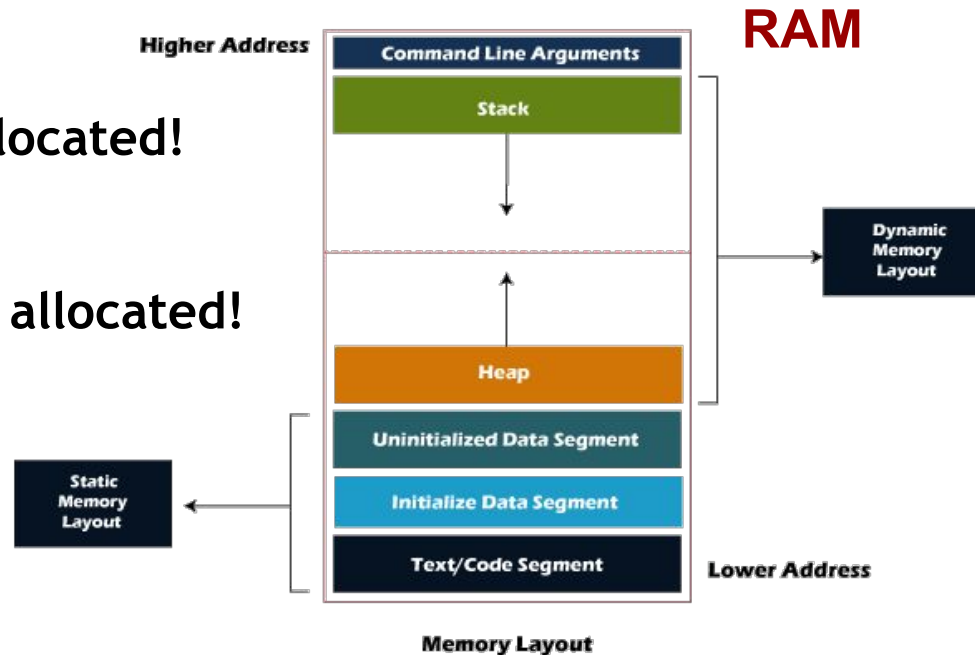
```
    char input[128];      ← 128 bytes allocated!
```

```
    char *ex = (char*) malloc(size);
```

```
    ...
```

```
}
```

↑
**'size' bytes
allocated!**



Why Use Heap?

Sometimes we don't know how much memory spaces we need until runtime!

For Example:

For Assignment 2 (the shell assignment), we have a constraint where there will be at most 5 processes running concurrently.

What if we remove that constraint? How do we keep track of all the processes? Could be 5, 10, 20, 30, 31, ...

Use a linked list where each linked list node lives on the heap!

Assignment 4: Implementing Your Own ‘malloc’

Memory Allocator is what manages the heap during program runtime to allow for dynamic allocation.

However, it presents a new set of challenges and design space...

For the purpose of Assignment 4, here are the terms you should know:

- **Block Splitting**
- **Block Coalescing**
- **Implicit Free List**
- **First-Fit Strategy**

But first, the design of Assignment 4's heap...

Assignment 4: The Heap Design

Your allocator is managing a “virtualized” heap.

Your “virtualized” heap lives in the .data segment (declared as a global variable):

```
unsigned char heap[64];
```

Think of indexes into “heap” as memory pointers.

Assignment 4: Heap Organized As Implicit Free List

You identify free block by traversing the whole heap. No explicit data structure keeping track of the free blocks (in contrast with *explicit free list*).

Use **first-fit strategy** to allocate the block: allocate the first free block that fits the requested size.

How to traverse the heap:

```
void allocateBlock(int sz){
    int i = 0; // pointer to beginning of heap
    int found = 0; // a flag
    while (!found && i < 64){
        if (blockIsFree(heap[i]) && (getSize(heap[i])-2)>= sz)
            // block is free and fits the requested size
            found = 1; break;
        i += getSize(heap[i]); // go to header of next block
    }
}
```

Assignment 4: Header and Footer Structure

Header and footer are one byte each. The one byte contains two information: block size (most significant 7 bits) and allocation status (least significant bit):

Supposed a block with size 20 (header + payload + footer) and an allocated status, how is it represented by its header and footer?

*Representation
Of The Byte In
Binary:*

Decimal 20 == 0x14 == 0b10100

0 0 1 0 1 0 0 1

Block Size

Allocation Status

Assignment 4: Header and Footer Structure

Use bit-masking and/or bitwise shifts to isolate either the block size or the allocation status!

```
// to get block size
// - bitwise AND with 0b 1111 1110
// - bitwise shift to remove allocation status
(header & 0xFE) >> 1
```

*Representation
Of The Byte In
Binary:*

Decimal 20 == 0x14 == 0b10100

0 0 1 0 1 0 0 1

Block Size

Allocation Status

Assignment 4: Block Splitting

Once you find a free block, you need to split the block if free block size minus allocated block size is greater than 2:

- allocated block size = header + payload + footer
- header + footer takes 2 bytes. *If after splitting, there's only 2 bytes of free space left, then there is no point splitting since there is no space for payload.*

Assignment 4: Block Splitting Example

Your Heap Initially:

Header

Block Size: 64, Allocation Status: 0

Payload

Footer

Block Size: 64, Allocation Status: 0

Assignment 4: Block Splitting Example

After malloc(20):

Note: blocks (allocated or not) are contiguous!

Block Size: 22, Allocation Status: 1
Block Size: 22, Allocation Status: 1
Block Size: 42, Allocation Status: 0
Block Size: 42, Allocation Status: 0

Assignment 4: Block Coalescing

WHY: reduce **external fragmentation**. You need to do forward and backward coalescing!

To determine if you need forward coalescing:

```
// find next block's header
int next_header = current_footer + 1
if (current_footer < 63 && blockIsFree(heap[next_header])){
    // perform forward coalescing
    // (1) change current_header's block size to account
    //     for next_header's block size
    // (2) update next_footer's block size. Current_footer
    //     now becomes part of the payload
    // (3) zero'd out each byte of payload. NOTE: this is an
    //     requirement for this assignment. Real malloc will
    //     not do this because of speed!
}
```

Assignment 4: Block Coalescing

WHY: reduce **external fragmentation**. You need to do forward and backward coalescing!

To determine if you need backward coalescing:

```
// find previous block's footer
int previous_footer = current_header - 1
if (current_header > 0 && blockIsFree(heap[previous_footer])){
    // perform backward coalescing
    // (1) update previous_header's block size. current_header
    //     now becomes part of the payload
    // go to previous header
    int previous_header = previous_footer + 1 - previous_block_size
    // (2) change current_footer's block size to account
    //     for previous_header's block size
    // (3) zero'd out each byte of payload. NOTE: this is an
    //     requirement for this assignment. Real malloc will
    //     not do this because of speed!
}
```

Questions?